Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

Digital-AV Software Development Kit provides the foundation for a fully working bible application), with no external dependencies. In fact, implementation with fewer than 1000 lines of code is possible, as demonstrated by the golang sources in this SDK. The SDK provides everything, including data and index files. Some developers have discovered that they can be up and running in under an hour. Easily jumpstart your development by working with the provided golang sources, or go all in from scratch with the programming language of your choice.

BREAKING NEWS: Direct support for Rust, C++, and FlatBuffers is debuting in the Z31 revision. See the addendums at the bottom of this document for status of each of these endeavors. I expect great progress in Q1 2023.

The base Digital-AV SDK (AVSDK) is entirely file based. There are zero dependencies and zero language bias (all programming languages can read files). File formats defined in this document use a consistent naming convention: the extent of each data file reveals the content and record type. The table to the right, defines the various extents of files that compose the SDK.

File Extent	File Type	Record Type	Content Type
*.dx	data	fixed length	binary
*.ix	index	fixed length	binary
*.dxi	data + index	variable length	binary
*.bom	MD5 checksums	newline delimited	text
*.ascii	informational	newline delimited	text

File extents identify the format/file type. Fixed-width data files have *.dx file extents, while fixed-width index files have *.ix file extents. A third file-type combines the data and index into a single file (and those binary files are always variable length). The *.bom file contains MD5 hashes which can be used at runtime to verify the file conforms to the release. The *.ascii files provide additional information for the developer and are not expected to be deployed to the end-user.

AV-Writ.dx file has three variants to handle disparate memory and/or system constraints. While the file formats are detailed later in this document, the table to the right is provided to summarize the records widths for each variant. (only a single AV-Writ* file need be deployed with your application). It is up to the developer to weigh the footprint versus features in that decision

AV-Writ variants	Size in bits	Size in bytes
AV-Writ.dx	176 bits	22 bytes
AV-Writ-128.dx	12byte	16 bytes
AV-Writ-32.dx	32 bits	4 bytes

Digital AV SDK – Record layouts & File inventory
Z-series Edition / Revision: 3.1

Digital-AV – Detailed description of file layouts

The weightiest data files are those named AV-Writ*.dx; these files contain a stream of words for each verse of each chapter for each book. As these are not text files, the records are compact, with some integer fields being an index lookup into another SDK file. In essence, the entire set of binary files implement an efficient database of word embeddings, designed to be compactly manifested in RAM. The AV-Writ*.dx files with the widest records are obviously also the most feature rich.

AV-Writ.dx (22 bytes per record)

Record # 0 bits	Hebrew | Greek 4 x uint16	Verse uint16	Caps 2 bits	WordKey 14 bits	Punc byte	Transition byte	PN+WClass uint16	POS uint32	Lemma uint16
0	0x391C 0 0 0	0	0x8___	0x0015 (in)	0x00	0xE8	0x00E0	0x40080470	0x0015
1	0x391C 0 0 0	0	0x0___	0x0136 (the)	0x00	0x00	0x0D00	0x00000094	0x0136
2	0x391C 0 0 0	0	0x0___	0x24F9 (beginning)	0x00	0x00	0x4010	0x000001DC	0x24F9
…	<< Beginning of Genesis 1 depicted above >>								
0xBDDB9	0x25A0 0 0 0	30698	0x8___	0x0136 (the)	0x00	0xE0	0x0D00	0x00000094	0x0136
0xBDDBA	0x25A0 0 0 0	30698	0x8___	0x2CB2 (revelation)	0x00	0x00	0x4010	0x000001DC	0x2CB2
0xBDDBB	0x0978 0 0 0	30698	0x0___	0x001D (of)	0x00	0x00	0x0400	0x80004206	0x001D
…	<< Beginning of Revelation 1 depicted above >>								
0xC0C91	0x1460 0 0 0	31101	0x0___	0x015C (you)	0x00	0x00	0x20C0	0x00083BBD	0x015C
0xC0C92	0x0F74 0 0 0	31101	0x0___	0x0036 (all)	0xE0	0x04	0x0D00	0x00000004	0x0036
0xC0C93	0x0119 0 0 0	31101	0x8___	0x018A (amen)	0xE0	0xFC	0x8000	0x8000550E	0x018A
<< End of Revelation 22:21 depicted above >>									

AV-Writ.dx begins with *Greek & Hebrew* Strong's numbers in the Old & New Testament. Each English word can have up to four Strong's numbers associated with it. Strong's numbers are an integer representation of the original Hebrew/Greek words from which the English words were originally translated. While words in the Old Testament can have a maximum of four Strong's numbers associated with a single English word.

Hebrew | Greek encodings

Strongs #1	Strongs #2	Strongs #3	Strongs #4
1st Strongs #	2nd Strongs #	3rd Strongs #	4th Strongs #

Digital AV SDK – Record layouts & File inventory
Z-series Edition / Revision: 3.1

The New Testament can only have a maximum of three Strong's numbers associated with a single English word. This is characteristic of the KJV translation, but four slots are reserved even for the Greek to maintain a fixed record width across entire bible. For more information on Strongs numbers, refer to the Strong's Exhaustive Concordance for additional backround information. Also note that **Verse** is an inline index-pointer to the corresponding AV-Verse index.

Capitalization bits and WordKey

Description	Bit Pattern (Hex)
English Word	0x3FFF (mask for lexicon lookup)
1st Letter Cap	0x8000 (example: Lord)
All Letters	0x4000 (example: LORD)

The next sixteen bits can be thought of as two distinct fields: the first of those is **Caps**: these 2-bits identify whether to apply capitolization rules to the lexical word. 0x8___ means to capitolize the first letter of the word (e.g. Lord). 0x4___ means to capitolize all letters of the the word (e.g. LORD). Clearly, in English, the first letter of the first word of a sentence is capitolized, and these bits facilitate all such capitolization rules. When no bits are set, this indicates that the word should be represented exactly as it appears in the lexicon. The remaining 14-bits are called the **WordKey** (a lookup key for the AV-Lexicon). Incidentally, the lookup key is equally compatible with the Lexicons found in the older 2018 SDK (Revision #K817).

The next field is the **Punctuation** byte. Each word can have certain punctuation applied either as a prefix to the word, or alternatively as a suffix. An example of prefix punctuation is an open parenthesis. There are numerous examples of suffix punctuation, such as period, comma, or close parenthesis. The puncuation byte also has bits to represent italisized words in the text and even mark the words spoken by Jesus, which some bibles represent as red-colored text.

Punctuation Byte

Description	Bits
PUNC::clause	0xE0
PUNC::exclamatory	0x80
PUNC::interrogative	0xC0
PUNC::declarative	0xE0
PUNC::dash	0xA0
PUNC::semicolon	0x20
PUNC::comma	0x40
PUNC::colon	0x60
PUNC::possessive	0x10
PUNC::closeParen	0x0C
MODE::parenthetical	0x04
MODE::italics	0x02
MODE::Jesus	0x01

Digital AV SDK – Record layouts & File inventory
Z-series Edition / Revision: 3.1

REVISION IDENTIFIERS

Digital-AV SDK:	Z31c
SDK Document:	Z31c

Person/Number (4 bits)

Description	Left-Most Nibble
Person bits	0x3--- (0b--11)
Number bits	0xC--- (0b11--)
Indefinite	0x0--- (0b--00)
1st Person	0x1--- (0b--01)
2nd Person	0x2--- (0b--10)
3rd Person	0x3--- (0b--11)
Singular	0x4--- (0b01--)
Plural	0x8--- (0b10--)
WH*	0xC--- (0b00--)

Within AV-Writ.dx and in AV-Writ-128.dx, Person/Number (PN) is the left-most nibble of the WordClass field. PN applies to pronouns and verb casing. Whereas Modern English is not morpologically rich when it comes to verb case, Early Modern English was slightly richer with additional pronouns and verb cases for Second-Person-Singular and Third-Person-Singular, each distinct from the Early Modern Plural counterparts. The Digital-AV captures and preserves all these disctinct case markings. For instance, **thy** is second-person singular whereas Early Modern English **you** is always plural form of this pronoun. AV-SDK encodes the markings for both person and number using the binary representation depicted in the table to the left.

WordClass (12 bits)

NounOrPronoun	0x-03-
Noun	0x-01-
Noun: unknown gender	0x-010
Proper Noun	0x-03-
Pronoun	0x-02-
Pronoun: Neuter	0x-021
Pronoun: Masculine	0x-022
Pronoun: Non-feminine*	0x-023
Pronoun: Feminine	0x-024
Pronoun/Noun: Genitive	0x-0-8
Pronoun: Nominative	0x-06-
Pronoun: Objective	0x-0A-
Pronoun: Reflexive	0x-0E-
Pronoun: no case/gender	0x-020
Verb	0x-1--
to	0x-200
Preposition	0x-400
Interjection	0x-800
Adjective	0x-A00
Numeric	0x-B00
Conjunction	0x-C0-
Determiner	0x-D0-
Particle	0x-E00
Adverb	0x-F00

Transition bits are a composition of Verse-Transitions and Segment-Markers. These represent a compact mechanism for data file traversal, obviating the need for leveraging additional index files. The five left-most bits mark book, chapter, and verse transitions. The three right-most bits mark linguistic segmentation [sentence and/or phrase] boundaries. In this edition of the SDK, these bondaries are interpretted based upon a combination of verse transitions and punctuation.

Verse Transitions

Description	5-bits
EndBit	0x10
BeginningOfVerse	0x20
EndOfVerse	0x30
BeginningOfChapter	0x60
EndOfChapter	0x70
BeginningOfBook	0xE0
EndOfBook	0xF0
BeginningOfBible	0xE8
EndOfBible	0xF8

> * *his* is used ambiguously in the Authorized Version for third-person-singular pronouns. *his* is either masculine or neuter (*its* appears just once in the sacred text). Therefore, *his* can neither be uniformly marked as masculine, nor neuter. Instead, we mark the genitive pronoun *his* as non-feminine.

Digital AV SDK – Record layouts & File inventory
Z-series Edition / Revision: 3.1

There are two additional trimmed down versions of the AV-Writ files which contain subsets of AV-Writ.dx. These can be used for more memory constrained implementations or utilized where the additional data fields are not needed.

AV-Writ-128.dx (16 bytes per record)

Record 0 bits	Hebrew | Greek 4 x uint16	Verse uint16	Caps 2 bits	WordKey 14 bits	Punc byte	Transition byte	PN|WordClass uint16
0	0x391C 0x0 0x0 0x0	0x0000	0x8___	0x0015 (in)	0x00	0xEF	0x00E0
1	0x391C 0x0 0x0 0x0	0x0000	0x0___	0x0136 (the)	0x00	0x00	0x0D00
2	0x391C 0x0 0x0 0x0	0x0000	0x0___	0x24F9 (beginning)	0x00	0x00	0x4010
…	<< Beginning of Genesis 1 depicted above >>						
C0C91	0x1460 0x0 0x0 0x0	0x797D	0x0___	0x015C (you)	0x00	0x00	0x20C0
C0C92	0x0F74 0x0 0x0 0x0	0x797D	0x0___	0x0036 (all)	0xE0	0x06	0x0D00
C0C93	0x0119 0x0 0x0 0x0	0x797D	0x8___	0x018A (amen)	0xE0	0xFE	0x8000
<< End of Revelation 22:21 depicted above >>							

AV-Writ-32.dx (4 bytes per record)

Record 0 bits	Caps 2 bits	WordKey 14 bits	Punc byte	Transitions byte
0	0x8___	0x0015 (in)	0x00	0xEF
1	0x0___	0x0136 (the)	0x00	0x00
2	0x0___	0x24F9 (beginning)	0x00	0x00
…	<< Beginning of Genesis 1 depicted above >>			
C0C91	0x0___	0x015C (you)	0x00	0x00
C0C92	0x0___	0x0036 (all)	0xE0	0x06
C0C93	0x8___	0x018A (amen)	0xE0	0xFE
<< End of Revelation 22:21 depicted above >>				

Segment Markers

Description	3-bits
HardSegmentEnd	0x04
CoreSegmentEnd	0x02
SoftSegmentEnd	0x01
RealSegmentEnd	0x06

Hard Segments: . ? !
Core Segments: :
Real Segments: . ? ! :
Soft Segments: , ; ( ) --

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

AV-Book index provides indicies into AV-Chapter, AV-Verse, and AV-Writ, and corresponding chapter-counts, verse-counts, and word-counts (for each of the sixty-six books of the bible). It reserves a fixed sixteen bytes for the book-name, a fixed nine bytes (2+3+4) for 2-character, 3-character, and 4-character abbreviations. The remaining nine bytes are a comma-delimited list of any additional alternate abbreviations.

AV-Book.ix (50 bytes)

Record # *0 bits*	Book Number *byte*	Chapter Count *byte*	Chapter Index *uint16*	Verse Count *uint16*	Verse Index *uint16*	Writ Count *uint32*	Writ Index *uint32*	Book Name *16 bytes* (utf8)	Abbreviations (utf8) a2 a3 a4 **(9 bytes)**	Alternates **(9 bytes)**
0	0	0	0	0	0	0	0x310c	Z31c-----------	-- --- ----	Revision-
1	1	50	0	1533	0	38262	0	Genesis---------	Ge Gen Gen-	Gn-------
2	2	40	50	1213	1533	32685	38262	Exodus----------	Ex Exo Exod	---------
3	3	27	90	859	2746	24541	70947	Leviticus-------	Le Lev Lev-	Lv-------
…										
66	66	22	1167	404	30698	11995	777656	Revelation------	Re Rev ----	

The dashes (-) above represent zero ('\0'). In the "a2 a3 a4" column, these nine bytes contain a 2-character abbreviation, a 3-character abbreviation, and a 4-character abbreviation, The AV-Book.ix file has an updated format, beginning with the Z31 revision. Note that the newer format now contains 67 records instead of 66. The zeroth record contains metadata about the revision and makes record #1 correspond to book #1. The previous AV-Book.ix has been renamed AV-Book-Z14.ix. Either file is fully compatible with all Z-series revisions. Weighing in with only 1,238 additional bytes, the newer file likely results in fewer ancillary lookups.

AV-Book-Z14.ix (32 bytes) *Compatible with earlier SDK releases, including Z14*

Record # 0 bits	Book Number byte	Chapter Count byte	Chapter Index Uint16	Book Name 16 bytes (utf8)	Book Abbreviations 12 bytes (utf8)
0	1	50	0	Genesis	Ge
1	2	40	50	Exodus	Ex
2	3	27	90	Leviticus	Le
…					
65	66	22	1167	Revelation	Re

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

AV-Chapter.ix (10 bytes)

Record # 0 bits	Writ Index uint32	Writ Count uint16	Verse Index uint16	Verse Count uint16
0x000 (genesis:1)	0	797	0	31
0x001 (genesis:2)	797	632	31	25
0x002 (genesis:3)	1429	695	56	24
. . .				
0x4A2 (revelation:20)	787852	477	31039	15
0x4A3 (revelation:21)	788329	749	31054	27
0x4A4 (revelation:22)	789078	573	31081	21

AV-Chapter-Z14.ix (8 bytes)

Record # 0 bits	Writ Index uint32	Verse Index uint16	Word Count uint16
0x000 (genesis:1)	0	0	797
0x001 (genesis:2)	797	31	632
0x002 (genesis:3)	1429	56	695
. . .			
0x4A2 (revelation:20)	787852	31039	477
0x4A3 (revelation:21)	788329	31054	749
0x4A4 (revelation:22)	789078	31081	573

NOTE:

AV-Chapter.ix differs from earlier revisions, as it how includes verse count and an altered column order than the earlier Z14 Revision.

AV-Chapter contains one extra data field than its Z14 counterpart, but both files are fully compatible with all Z-series revisions.

AV-Verse.ix (4 bytes)

Record# 0 bytes	Book, Chapter, Verse, Words 4 bytes: BB:CC:VV:WordCnt	
0x0000	1:1:1:10	In the beginning …
0x0001	1:1:2:29	And the Earth …
0x0002	1:1:3:11	And God said …
...		
0x797B	66:22:19:44	And if any man … are written in this book.
0x797C	66:22:20:16	He which testifieth … Even so, come, Lord Jesus.
0x797D	66:22:21:12	The grace of our Lord … be with you all. Amen

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

The AV-Lemma file originally appeared in the 2017 Edition of the SDK. The original version obtained Lemmata from the NLTK Python library. Now Lemmata are obtained from the MorphAdorner Java server (MorphAdorner also performs all of the POS tagging). Incidentally, each Lemma ordinarily maps to multiple English words or lexemes, (e.g. 'be' is the lemma of 'are', 'were', 'is', 'art', 'wast', and 'be'). Moreover, words like 'run' can function both as a verb and a noun.

AV-Lemma.dxi (variable length)

Part-of-Speech uint32	Word Key uint16	Word Class uint16	Count uint16	Lemma Array Uint16[] (Word or OOV keys)
0x00000036	0x0001 (a)	0x0F00	1	0x0001
0x00000094	0x0001 (a)	0x0D00	1	0x0001
0x80004206	0x0001 (a)	0x0400	1	0x0001
0x01074F9C	0x0002 (i)	0x4080	1	0x0002
…				
0x00003A1C	0x027A (elim)	0x4030	1	0x027A
0x000001DD	0x027B (elms)	0x8010	1	0x8304 (OOV: elm)
…				
0xFFFFFFFF				

Part-of-Speech needs to be considered when accessing the lemma utilizing AV-Lemma for looking up the lemma for a word. Lemmas contains a list of WordKeys and/or OOVKeys. When a Lemma is OOV[1], it cannot be found in the Lexicon, but it can be found in the OOV table.

AV-Lemma-OOV.dxi (lookup for OOV lemmas)

OOV Key uint16	OOV Word Length+1 bytes
0x8301	aid\0
…	
0x8F01	covenantbreaker\0
…	

OOV (composition by example)

OOV Marker 1 bits	OOV Length 7 bits	OOV Index byte
0x8___	0x_7__	0x__01

(binary of 0x8301 is b1000001100000001)

---

[1] OOV stands for "Out of Vocabulary": Not all lemmas are in the AV-Lexicon; these OOV words can be looked up in the AV-Lemma-OOV table. As an example, "covenantbreakers" is in the KJV bible and therefore in the lexicon. However, covenantbreaker is not in the KJV bible (It is an example of an OOV word).

Digital AV SDK – Record layouts & File inventory
Z-series Edition / Revision: 3.1

AV-Lexicon provides both original and modern orthographic representations for each lexeme identified in AV-Writ. It also includes a search-version of the lexeme that strips out all hyphens. Next, an array each Part-of-Speech (POS) associated with that lexical item is identified by a encoded value using 5-bit character encoding. A reference implementation of to decode a Uint32 value into a human readable POS string can be found in the github repo.

AV-Lexicon.dxi (data and index combined: variable length records)

Rec # (0 bits)	Entities uint16	Size uint16	POS[0] uint32	POS[1] uint32	POS[2] uint32	•••	POS[n-1] uint32	Search char [ ]	Display char[ ]	Modern char [ ]	
0	0xFFFF	n=2	12567	0x310C							metadata
1	0x0000	n=4	0x00000094	0x00000036	0x0000000A		0x80004206	a			Entities = { } dt, av, j, pp-f
2	0x0000	n=3	0x01074F9C	0x0000000A	0x01073F9C			i			Entities = { } pns11, j, pno11
3	0x0000	n=1	0x000002A8					o		oh	{ } oh
…											
366	0x8009	n=2	0x00003A1C	0x000740FC				adam			Entities = {Man, City} np1, npg1
…											
1311	0x0000	n=2	0x01073FBC	0x0000000A				thou		you	Entities = { } pns21, j
…											
12567	0x0000	n=1	0x0000000A					Mahershalalhashbaz	Maher-shalal-hash-baz		Entities = { } j

Entities = {Hitchcock=0x8000, men=0x1, women=0x2, tribes=0x4, cities=0x8, rivers=0x10, mountains=0x20, animals=0x40, gemstones=0x80, measurements=0x100}

NOTE: AV-Lexicon.Z31 differs from Z14 revision: it inserts a zeroth-record, making lex-key equal to record-index. It also differs by omitting the marker/final record after record #12567, as did the Z14 Revision. Otherwise, they are identical.

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

AV-WordClass.dxi (data and index combined: variable length records)

WordClass uint16	Width uint16	1st POS uint32	2nd POS uint32	3rd POS uint32	. . .	Nth POS uint32	
0x0010	N=4	0x4000394E	0x00003950	0x40075AC7		0x40075ACE	n-jn, njp, n-vvg, n-vvn
0x00E0	N=1	0x01074F9C					p-acp
0x0100	N=29	0x00005842	0x000B0893	0x00005847		0x00005ADA	Vbb, vbds, vbg, … vvz
. . .							

It should be noted that both the 16-bit WordClass field and each POS field contains part-of-speech information, but the 16-bit WordClass field is more granular and has a bitwise representation. Contrariwise, the encoded 32-bit POS fields have far more fidelity, but require decoding.

For more information, see:

- https://github.com/kwonus/Digital-AV/blob/master/z-series/Part-of-Speech-for-Digital-AV.pdf
- https://github.com/kwonus/AVXText/blob/master/FiveBitEncoding.cs [*method signature:* **string DecodePOS(Uint32 encoding)**]

AV-Names.dxi (data and index combined: variable length records)

WordKey uint16	1st Meaning	Delimiter	2nd Meaning	Delimiter	3rd Meaning	Delimiter	• • •
AVLexicon WordKey for Aaron	a teacher	|	lofty	|	mountain of…	\0	
AVLexicon WordKey for Abaddon	the destroyer	\0					
AVLexicon WordKey for Abagtha	father of the…	\0					
. . .							

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

avx.go (golang source code)

avx.go implements a web-server (HTTP server) that provides the entire text of the AV bible utilizing AVX extensions, but still uses simple semantics. Version numbers for source code are respective of the SDK Document revision numbers. The first release of avx.go, which had been updated to the z-series SDK, was the Z081 golang source-code revision.

> *Example of GoLang source in operation may be available at avbible.net:*
>
> https://avbible.net/avx/
> (the web-site above also utilizes NGINX as a reverse-proxy for HTTPS)

There are a couple of URLs used for testing and validation. They also illustrate how avx.go can be extended:

- http://localhost:2121/
- http://localhost:2121/help
- http://localhost:2121/validate

The / endpoint simply reports the release number of the optional avx.go web-server component. The /help endpoint provides primitive information about the web-service. /help can be easily replaced by

> *NOTES:*
>
> 1. As the web-server is not hardened, it should be placed behind a reverse-proxy if exposed to the open Internet. This is a common pattern; Apache httpd, NGINX, Caddy, or IIS can easily be configured to serve as a reverse-proxy.
> 2. URL form #3 and #5 are discussed under the description of the *.avspec format

developer. The /validate endpoint reports on the validity of data files in accordance with the bom (The "bom", or bill of materials, is described in the section labelled AV-Inventory.bom later in this document). In addition to the administrative URL's described above, here is a list of the foundational endpoints that provide the core functionality of avx.go:

1. http://localhost:2121/avx/genesis
2. http://localhost:2121/avx/genesis/1
3. http://localhost:2121/avx/gen/1?sessionID
4. http://localhost:2121/avx/rev/22?sessionID=day&amen
5. http://localhost:2121/avx/rev/22?sessionID=$FFFFFFFFFFFF
6. http://localhost:2121/avx/css/sessionID.css

All of these endpoints can be summarized as one of two types: getting the chapter of a book, or getting a CSS stylesheet. When no chapter is provided, chapter 1 is always implied. When no session identifier is provided, the resulting chapter request is decorated with the baseline stylesheet, named /css/AV-Stylesheet.css. When a session identifier is provided, the session number dictates the name of the CSS file that will decorate the chapter request. Moreover, avx.go can compile information into a CSS stylesheet. When a request is made for Genesis using the URL depicted in #3 above, a stylesheet becomes linked in the response to a stylesheet with the URL depicted in #8 above. A web-browser will make an immediate subsequent request to get the stylesheet. If /css/sessionID.css does not exists, avx.go will automatically compile a file named /css/sessionID.avspec. Similarly, but easier to understand in #4 above, the URL would generate CSS which would highlight the words **day** and **amen**. In order to maintain optimal performance, session identifiers are non-volatile. In order to overwrite a *.css files and/or *.avspec files, they must be manually deleted beforehand. Avx.go uses Z08 edition.

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

*.avspec file format

WordKey Count Uint16	Array of Uint16	
0xnnnn	0xnnnn Count of WordKeys is followed by WordKey list [corresponds to AV-Lexicon]	
BookChapter Uint16	**Verse Count byte** *(matching verses)*	**Array of byte**
0xbbcc	0xkk	0xkk Count of matching verses is followed by an array of Verse numbers
…		
0xbbcc	0xjj	0xjj Count of matching verses is followed by an array of Verse numbers
0x0000		

avx.go software ignores everything after the first record above. Only that first record defines the CSS file. And that first line is expanded word-for-word into highlights for each supplied wordkey. A slight variation here is that Strong's numbers will eventually also support highlighting. To highlight Strong's numbers, set the 0x8000 bit for Hebrew and the 0x4000 bit for Greek. The URL form that was depicted with this syntax, sessionID=$FFFFFFFFFFFF, is primarily intended for testing. Here, the hex digits that follow the dollar sign ($) are expected to be expansions of the format described above (No record separators, just a representation of the raw bytes described above, in Big-Endian order).

AV-Stylesheet.css (text file containing CSS for avx.go; optional)

This standard-format CSS stylesheet should be included when avx.go is utilized in your development. This optional stylesheet is included in the SDK, but it can be customized in any way by the web designer. However, the web designer should realize that any references in the CSS to image files will result in 404 errors unless support is explicitly added to avx.go by your development team. Finally, avx.go always links chapter output to the AV-Stylesheet.css stylesheet, even when a *.avspec derived stylesheet is also specified.

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

AV-Inventory-Z31.bom (text file which identifies core inventory)

This is an ascii text file that provides a bill-of-materials for the delivered files. For each file of the SDK, the bom contains a line item for the artifact. Each line has three fields, separated by whitespace:

1. Name of the file
2. MD5 hash (hexadecimal representation) of the file
3. Record Length (decimal representation // uint32) of the file // 0 for variable-width files
4. Record Count (decimal representation // uint32) of the file
5. Size in bytes (decimal representation // uint32) of the file

The avx.go server implements a validation function, using an older bom; it reads the bom and reports inconsistencies. To avoid malicious attack, utilization of the bom is highly recommended, but not required. It helps mitigate corruption, both intentional and unintentional.

AV-Inventory-Z31.md5 (new with the Z31 release)

This ascii text file contains the MD5 of AV-Inventory-Z31.bom. For utmost security, utilized this MD5 to check the validity of AV-Inventory-Z31.bom, in addition to checking the validity of each file utilized at runtime from the bom.

OVERALL PROJECT STATUS:

It's an exciting time at AV Text Ministries, and if you want to lend a hand. Let us know your technical skills and interests and we can help jumpstart you onto the team. We are embarking on brand-new support for Rust. Currently, AV Text Ministries is 100% volunteer, so if you don't just have passion about the mission as your raw motivation, it might not be the best fit.

Finally, on the non-technical side of things, we would certainly welcome a ministry sponsor that would want to place AV Text Ministries under the banner of their own local church ministry. Check http://avtext.org to discover our overall vision.

HOW THE DIGITAL-AV "PLATES" ARE AUTHORED:

Initially, various publicly available KJV texts were parsed and dutifully compared (comparing scripture with scripture [1 Corinthians 2:13]). That work produced the freeware program, AV-1995 for Windows; it was written in Delphi/Pascal and was maintained until the AV-2011. In 2008, the initial Digital-AV SDK was conceived and produced, harvesting much of the inner workings of AV-2008, utilizing RemObjects Oxygene/Pascal as a development platform and releasing it as open source. Later, AV-2011 was "compiled" using AV-2008 as a baseline. Subsequently, the 2017/2018 Editions were "compiled" using AV-2011 as a baseline. The Z07 revision of the SDK were baselined from AV-2018 edition using the K817 revision. C# is now the programming language of the SDK compiler; and the ancient pascal sources were finally retired (replaced by C# sources) in 2018. The SDK-compiler uses MorpAdorner[2] (written in Java 1.6) and the NUPOS [3]tag-set. NLTK[4] (Python) used when MorphAdorner encounters a word out of its vocabulary. Java and Python dependencies are not exhibited in the delivered SDK (They are only part of the compilation process).

---

[2] http://morphadorner.northwestern.edu/morphadorner/
[3] https://github.com/kwonus/Digital-AV/blob/master/z-series/Part-of-Speech-for-Digital-AV.pdf
[4] http://www.nltk.org

Digital AV SDK – Record layouts & File inventory

Z-series Edition / Revision: 3.1

LICENSE REQUIREMENT:

- In order to comply with the MIT-style open-source license, please include AV-License.txt with your distribution of any file identified in this SDK. The text of that file, as of 2021, is provided also at the bottom of this page.

All SDK artifacts are on github.com:

https://github.com/kwonus/Digital-AV

IMPROVEMENTS & CAVEATS:

- Fundamental SDK format has stabilized and is substantially similar to the 2017 and 2018 editions.
- Another two fields have been added to AV-Writ which provides a precise Part-of-Speech representation and lemmatization for each word.
- AV-Lemma has also been updated in the 2020 SDK. Moreover, the binary format of AV-Lexicon is also substantially different from earlier editions.
- Part-of-speech (POS) bits were introduced into the SDK with the HA29 release. As of the Z07 release, POS bits have been substantially revised as the SDK now uses MorphAdorner for part-of-speech marking instead of NLTK (NLTK doesn't recognize archaic verbs and pronouns, whereas MorphAdorner does to some degree).
- The sqlite lexicon has been eliminated from the SDK.

ADDITIONAL RELEASE NOTES:

#1 The "Z-series" edition of the SDK introduced an updated revision number from earlier editions. Digital-AV revision numbers now use a three-digit character sequence, plus an optional suffix/subscript. All revision numbers now begin with the letter **Z**. The next two characters represent year and month of the revision. The character sequence is **Zym** where the first letter is always **Z**, indicating that this is the "Z-series" edition of the SDK (distinguishing it from older/legacy SDK editions); *y* represents the year, and *m* represents the month of the release. *y* encodes the year as a single base-36 digit; For example, (y=0) represents 2020; (y=1) represents 2021; (y = A) represents 2030; (y = K) represents 2040; (y = U) represents 2050. With respect to months, digits 1 through 9 are as expected; (m = A) is October; (m = B) is November; and (m = C) is December. An optional one-digit suffix/subscript may also be used. If the subscript is a Greek letter (α or β), then this identifies an alpha or beta release of the SDK. Otherwise, a suffix/subscript identifies the discrete date of the release, encoded in base-32; the 1st is 1, the 31st is v.

#2 Two revision numbers exist: Digital-AV SDK revision (aka, the "plate" revision) is the most significant set of files. Not all files in this SDK are required to produce working bible software. Incidentally, the sample source code provided in avx.go implements a minimal set of SDK artifacts, while still providing access to the entire AV Bible text. The avx.go sources use a slightly older version of the SDK (Z08 Revision) and would require minor adaption to update to this latest revision.

#3 Many of the binary files also have corresponding text files with an .ascii extent. These files are not provided for runtime execution. Instead, they should be considered as ancillary documentation to shed light, in painstaking detail, on the corresponding binary files.

#4 The Z31 SDK release is substantially identical/compatible with the previous Z14 release, the SDK now includes an addendum for Flat Buffers IDL and binary files

#5 The Z31 SDK release adds direct support for Rust, C, and FlatBuffers. Each addendum describes the specialized support and overall status. Code-Generation is utilized to extend support. These addendums are supplied by the dotnet project named SerializedFromSDK.csproj in the Z-Series/FB folder within the github repo.

Digital AV – Addendum for FlatBuffers (FB)

Z-series Edition / Revision: 3.1

If the developer is willing to take on the dependency of FlatBuffers[5], the deserialization can be driven using the IDL provided in the FlatBuffers Schema (fbs) IDL files. All binary-content files for deserialization for FlatBuffers have an ".data" extent. Corresponding IDL files have an ".fbs" extent. The layouts are substantially similar to the baseline SDK. Therefore, the baseline SDK documentation can still be consulted. However, the IDL (i.e. *.fbs file) can drive deserialization through the FlatBuffers code-generated sources for most programming languages.

The files in the table below are consistent with the latest revision of the baseline SDK. The fundamental difference is the serialization format itself.

Baseline AV SDK item	Baseline Size	Flatbuffer IDL	FlatBuffer binary content	FlatBuffer Size
AV-Writ.dx	17 mb	Written.fbs	Written.data	18 mb
AV-Book.ix	3 kb	Book-Index.fbs	Book-Index.data	1 kb
AV-Chapter.ix	12 kb	Chapter-Index.fbs	Chapter-Index.data	1 kb
AV-Verse.ix	122 kb	Verse-Index.fbs	Verse-Index.data	122 kb
AV-Lemma.dxi	179 kb	Lemmata.fbs	Lemmata.data	742 kb
AV-Lemma-OOV.dxi	8 kb	Lemmata-OOV.fbs	Lemmata-OOV.data	44 kb
AV-Lexicon.dxi	241 kb	Lexicon.fbs	Lexicon.data	1.2 mb
AV-WordClass.dxi	1 kb	WordClasses.fbs	WordClasses.data	3 kb
AV-Names.dxi	60 kb	Names.fbs	Names.data	292 kb

The FlatBuffers-special files identified above can be found in the FB sub-folder of the Z-Series SDK[6].   These files have been written using FlatSharp[7]. As of the date of this documentation, all FlatBuffer assets should be considered Alpha-quality. They are available for use, but completely untested as yet.

The status of FB support is entirely experimental and currently moth-balled until interest develops.

---

[5] See https://google.github.io/flatbuffers/
[6] See https://github.com/kwonus/Digital-AV/tree/master/z-series/FB
[7] See https://github.com/jamescourtney/FlatSharp

Digital AV – Addendum for Rust

Z-series Edition / Revision: 3.1

In general, the Rust generated files in Digital-AV/z-series/foundations/rust/ folder on github, define structures in lockstep with the binary files of the SDK. However, one major deviation is that the AV-Writ.dx file is segmented into 66 different structures (one for each book of the bible).

There are other minor deviations that should be intuitive by examining the struct definitions. These are driven somewhat by the rules of Rust. The value of the generated files is that no deserialization operations are required. Moreover, the implementation uses Rust arrays with static initializations.

These sources are in flux and undergoing active development. If you encounter problems utilizing these sources, please let me know.

Digital AV – Addendum for C/C++

Z-series Edition / Revision: 3.1

In general, the C/C++ generated files in Digital-AV/z-series/foundations/cpp/ folder on github, define structures in lockstep with the binary files of the SDK. There are some minor deviations that should be intuitive by examining the struct definitions. The value of the generated files is that no deserialization operations are required. The SDK is manifest using static C++ initializers.

The C++ generated sources are ahead of the Rust sources with respect to compilability. If you encounter problems, I would be excited to directly engage with you.